Lecture 26 — Timestamp, Validation, & Multiversion Protocols

Jeff Zarnett

`jzarnett@uwaterloo.ca`

Department of Electrical and Computer Engineering
University of Waterloo

March 26, 2018

We can also use timestamps for arranging concurrency.

The timestamp, as you will recall, is a unique identifier, and they are typically assigned in the order the transactions are submitted to the system.

If timestamps are used appropriately, then consistent results are produced without the use of any locks or locking, meaning deadlock can never occur.

The notation for the timestamp of a transaction T is $TS(T)$.

The transactions can be ordered based on their timestamps.

If a transaction T_i arrives and then later another transaction T_j arrives, then $TS(T_i) < TS(T_j)$.

If two transactions arrive at exactly the same time, then some sort of serialization procedure will be needed to make sure that they are not identical.

It is commonly the case that the system clock is used to get the time stamp; just read the current value of the clock and use that.

The alternative is a logical counter: increment a simple integer counter for each transaction.

After a timestamp is assigned, just increment the counter.

Eventually there is a limit (even if you can have 2^{64} transactions, in theory, which is quite a lot, but not infinite).

At some point the transaction counter will need to roll over or reset.

Generation of the timestamps is simple, it would seem, but what are they for?

Timestamps are used for serializability order in the schedule.

The system must ensure that the schedule executed is equivalent to a serial schedule in which transactions are ordered according to their timestamps.

Read and Write Timestamps

Every data element in the database is associated with separate timestamps for reading and writing.

The notation can vary: they are called **read_TS** and **write_TS** or **r-timestamp** and **w-timestamp**.

To save space I might personally like TS_r and TS_w.

Each time the a data element X is read or written, its appropriate timestamp is updated to the timestamp of the transaction performing the read or write.

The formal definition of the read timestamp of an item X: the largest timestamp amongst all the timestamps of transactions that have successfully read X.

So $TS(X) = TS(T_k)$ where T_k is the youngest transaction that has read X successfully.

The formal definition of the write timestamp of an item X: the largest timestamp amongst all the timestamps of transactions that have successfully written X.

So $TS(X) = TS(T_k)$ where T_k is the youngest transaction that has written X successfully.

Unfortunately, simple timestamp ordering does not avoid the risk of cascading rollback.

Schedules are not guaranteed to be recoverable.

This leads us to the timestamp ordering protocol.

For a read operation T_i is performing on X:

1. If $TS(T_i) < TS_w(X)$ then T_i needs to read a value of X that was already overwritten – and therefore the read operation is not permitted and T_i is rolled back.
2. If $TS(T_i) \geq TS_w(X)$ the read operation is executed and TS_r is set to $\max(TS(T_i),\ TS_r(X))$.

For a write operation T_i is performing on X:

1. If $TS(T_i) < TS_r(X)$ then the value that T_i is producing was needed previously but the system assumed it wasn't going to happen and proceeded without it; thus the write is not permitted and T_i is rolled back.
2. If $TS(T_i) < TS_w(X)$ then T_i is attempting to write an obsolete value; the write is not permitted and T_i is rolled back.
3. Otherwise, the write proceeds and $TS_w(X)$ is updated to $TS(T_i)$.

When a transaction is rolled back and restarted, the transaction is assigned a new timestamp.

The timestamp ordering protocol does provide us with conflict serializability.

Starvation can still occur: a very long transaction might be constantly rolled back by shorter transactions constantly swooping in and making changes.

Unfortunately it might be necessary to block other transactions and let the long one finish.

In the example below, transaction T_{25} reads values A and B , and transaction T_{26} modified both A and B:

T_{25}	T_{26}
read(B)	
	read(B)
	$B := B - 50$
	write(B)
read(A)	
	read(A)
display($A + B$)	
	$A := A + 50$
	write(A)
	display($A + B$)

It is worth noting that neither two phase locking nor the simple timestamp ordering protocol covers all serializable schedules.

If we would like to recoverability we can have strict timestamp ordering.

T_i that issues a read or write of X where $TS(T_i) > TS_w(X)$, then the read or write is delayed until the transaction that last wrote X has committed or aborted.

This, in some way, requires simulating locking item X until a previous write has been committed or aborted.

There cannot be a deadlock because a transaction can only wait on an older transaction (one with a lower timestamp).

A modification of the basic algorithm called Thomas's Write Rule allows greater concurrency than the basic algorithm and also, hopefully, rejects fewer writes.

To explain it simply, it is "ignore outdated writes".

More formally, the write item checks are modified to the following;

1. If $TS_r(X) > TS(T_i)$ then abort and roll back T_i (write rejected).
2. If $TS_w(X) > TS(T_i)$, then do not execute the write, but continue executing (do not roll back the transaction).
3. Otherwise, the write proceeds and $TS_w(X)$ is updated to $TS(T_i)$.

A transaction that is read-only does not interfere with any other read-only transaction.

If a database has many more reads than writes, then a concurrency control routine might do unnecessary work or delay transactions unnecessarily.

If we could instead just step in where transactions may be in conflict, that would be better.

It is necessary then to find out what transactions are the ones that have the potential for conflict...

Transactions have three phases (although read only transactions skip the last one).

1 **Read Phase.**
2 **Validation Phase.**
3 **Write Phase.**

The validation test for the current transaction T_i, for each transaction T_k where $TS(T_k) < TS(T_i)$, one of the following conditions must hold:

1. T_k finishes before the start of T_i: serializability order is maintained because the current transaction starts after the previous one finished.

2. The set of data items written by T_k does not intersect with the set of data items read by T_i, and T_k finishes writing before T_i starts its validation phase. This ensures that the writes of these transactions do not overlap.

To make this work out as expected we need, now, three timestamps for each transaction.

One for the start of the transaction, one for the start of the validation phase, and then a final timestamp for when the transaction is finished.

Validation prevents cascading rollbacks because the writes don't really take place until the validation phase has taken place.

This is called optimistic concurrency control: transactions assume they will finish.

The locking and timestamp approaches re a bit more pessimistic. They force a rollback if there is a risk of conflict.

That potentially leads to an interesting discussion on the subject of whether an optimistic or pessimistic type approach is better.

The general recommendation is that optimistic approaches are lower overhead and better for situations where we expect there is not much contention.

T_{25}	T_{26}
read(B)	
	read(B)
	$B := B - 50$
	read(A)
	$A := A + 50$
read(A)	
< validate>	
display($A + B$)	
	< validate>
	write(B)
	write(A)

Thus far we have said that a transaction must be either delayed or denied (aborted) to ensure serializability.

In a multiversion scheme, we have multiple versions of the data at the same time.

Instead of rejecting a read because a write has taken place in the meantime, the read is provided an older version of the item.

Choosing which one is meant to be read is the hard part of this routine.

This does require more storage space to maintain multiple versions of the data.

Old versions might be maintained anyway for recovery, or just record keeping purposes.

Many databases are used in environments where all changes must be tracked for legal reasons.

Ensuring serializability requires holding to just two rules:

1. If transaction T tries to write item X, and version i of X has the highest write timestamp of all versions of X that is also less than or equal to the timestamp of T, and the read timestamp of X_i is larger than the timestamp of T, the abort the transaction T and roll it back. Otherwise, create a new version of X with read and write timestamps equal to the timestamp of T.

2. If transaction T wants to read item X, find the version of X (X_i) that has the highest write timestamp of all versions of X that is also less than or equal to the timestamp of T. Update the read timestamp of X_i to the maximum of its current read timestamp and the timestamp of T.

The goal is that read operations are always successful, even if they read an older version.

A write, however, might not be permitted if it comes "too late" – trying to write a version that another transaction would have read.

The drawback is that every read requires us to update the read timestamp, which is itself a write operation (and may result in a disk access).

When do we know that we can get rid of an old version?

An old version is no longer needed when its write timestamp is smaller than that of the oldest transaction in the system.

But one must be careful then to only delete old versions; at least the current version must be maintained.

But also it means that in a period in which there are no transactions active, we may clean up any old versions.

That may not occur frequently enough in a heavily used database, so it might be a periodic task to simply trash very old versions.

We resolve any conflicts on writes through rollbacks, which are expensive, rather than through waiting.

What if we said we want an all-of-the-above solution that combines two phase locking with the multiversion scheme?

The advantage of multiversion schemes is that reads don't fail.

The advantage of two phase locking is that locks allow transactions to wait rather than roll back, and deadlock is avoided.

The plan is to allow each transaction to be handled in the way they like best.

Thus, read-only transactions are differentiated from the update transactions.

Update transactions use two phase locking, with the restriction that they hold locks up until the end of the transaction.

There his only one timestamp for the transaction, which is used to both serialize the transactions and to assign the version of the data written.

The timestamp under this scheme is an incrementing counter rather than the actual time of the transaction.

Read-only transactions are assigned a timestamp based on the current value of the counter for update transactions, at the time when the read transaction is created.

Multiple read transactions can have the same value of their timestamp.

In effect, the read transaction's timestamp k is declaring that its view of the data is the version at time k.

A transaction that reads some data X with its timestamp k gets the version of X with the largest timestamp less than or equal to k.

If an update transaction wants to read an item, it acquires a shared lock on the latest version of the item.

To do an update it acquires an exclusive lock and creates a new version of that item with the timestamp being the maximum value (effectively: infinity)

The new version has this infinite timestamp so it won't be read by any read transaction, but it will be overwritten when the write transaction commits.

This does mean that only one transaction can commit at a time.