

# Tutorial 5 — Query Optimization, Planning, Evaluation

Richard Wong  
`rk2wong@edu.uwaterloo.ca`

Department of Electrical and Computer Engineering  
University of Waterloo

March 4, 2018

- 1 What are the two metrics we will use to estimate query operation costs?
- 2 What does each metric represent?
- 3 How do we use the metrics to arrive at an estimate?

- 1 In this course, we care primarily about cost contributed by block transfers and by disk seeks.
- 2 A block transfer occurs whenever we need to retrieve a block from disk to read in memory, or when we need to write a block back to disk.  
A disk seek occurs when we need to access a disk block non-sequentially.
- 3 Let block transfers take time  $t_{transfer}$  and disk seeks take time  $t_{seek}$ .  
Let there be  $b$  block transfers and  $s$  disk seeks.  
Then we estimate the cost of a query operation with the formula  $b(t_{transfer}) + s(t_{seek})$ .

Suppose we run a query that performs a single-attribute GREATER THAN comparison in its WHERE clause.

e.g.

```
SELECT * FROM people  
WHERE age > 20
```

How might the following evaluation strategies impact the (worst-case) cost of the operation?

- 1 Use a primary index if there is one.
- 2 Use a secondary index if there is one.
- 3 Use a linear scan.

What if we were performing a LESS THAN comparison?

Assuming our indices are B+ trees, let  $h_{index}$  be the height of the index.

Let  $B$  be the number of blocks in the relation.

Let  $b$  be the number of blocks containing records that satisfy the condition.

Let  $n$  be the number of records that satisfy the condition. ( $n \geq b$ )

**1** Primary index:  $h_{index}(t_{seek} + t_{transfer}) + b(t_{transfer})$

$h_{index}$  seeks and transfers to locate the starting point with the index.

Since this is a primary index, the last seek in the index brings us straight to the data, and no additional seek is required.

$b$  transfers to read the sequentially-organized data, following leaf node pointers.

**2** Secondary index:  $(h_{index} + n)(t_{seek} + t_{transfer})$

$h_{index}$  seeks and transfers to locate the starting point with the index.

Since this is a secondary index, we need a seek and a transfer for each block. In the worst case, each subsequent record is located on a different block.

$n$  seeks and transfers for those blocks.

**3** Linear scan:  $t_{seek} + B(t_{transfer})$

1 seek,  $B$  transfers to read the whole relation. No index  $\rightarrow$  no ordering assumption.

## Exercise 5-2 Solution, continued

If we were performing a LESS THAN as opposed to a GREATER THAN comparison, using the primary index will give no benefit, since we only need to seek to the beginning of the relation and perform a linear scan (assuming ascending order of index).

Note that using a secondary index could result in a greater cost than a linear scan, especially if  $n \gg b$ .

Suppose there are  $b_r$  blocks in  $R$  and  $b_s$  blocks in  $S$ .

Derive the worst-case and best-case cost estimate for a block nested-loop join,

$R \bowtie_{\theta} S$ .

In the worst case, we can assume that we have enough memory to hold exactly 1 block of  $R$  and exactly 1 block of  $S$  in main memory.

For each block in  $R$ , of which there are  $b_r$ , we have to:

- 1 **seek** to the block in  $R$ ,
- 2 **transfer** that block to main memory, then
- 3 **seek** to the beginning of  $S$ .
- 4 then for each block in  $S$ , we need to:
  - 1 **transfer** that block to main memory, then
  - 2 do some CPU work to accumulate the  $R, S$  record pairs that match the predicate.

So the join cost would be  $2b_r(t_{seek}) + b_r(1 + b_s)(t_{transfer})$ .



## Exercise 5-3 Solution, continued

In the best case, we can assume that we have enough memory to hold all of  $R$  and all of  $S$  in main memory, so we just need to:

- 1 **seek** to the beginning of  $R$ ,
- 2 **transfer** the  $b_r$  blocks of  $R$  to main memory,
- 3 **seek** to the beginning of  $S$ ,
- 4 **transfer** the  $b_s$  blocks of  $S$  to main memory, then
- 5 do some CPU work to accumulate the result.

So the join cost would be  $2(t_{seek}) + (b_r + b_s)(t_{transfer})$ .