

Tutorial 7 — Data Mining and Transactions

Richard Wong

`rk2wong@edu.uwaterloo.ca`

Department of Electrical and Computer Engineering
University of Waterloo

March 11, 2018

What is a data warehouse, and why is it useful to have one?

Exercise 7-2

Suppose we are trying to predict the value *Wait*.

Between attributes *Pat* and *Type*, which is better to split a decision node on?

Example	Attributes										Target
	<i>Alt</i>	<i>Bar</i>	<i>Fri</i>	<i>Hun</i>	<i>Pat</i>	<i>Price</i>	<i>Rain</i>	<i>Res</i>	<i>Type</i>	<i>Est</i>	<i>Wait</i>
X_1	T	F	F	T	Some	\$\$\$	F	T	French	0-10	T
X_2	T	F	F	T	Full	\$	F	F	Thai	30-60	F
X_3	F	T	F	F	Some	\$	F	F	Burger	0-10	T
X_4	T	F	T	T	Full	\$	F	F	Thai	10-30	T
X_5	T	F	T	F	Full	\$\$\$	F	T	French	>60	F
X_6	F	T	F	T	Some	\$\$	T	T	Italian	0-10	T
X_7	F	T	F	F	None	\$	T	F	Burger	0-10	F
X_8	F	F	F	T	Some	\$\$	T	T	Thai	0-10	T
X_9	F	T	T	F	Full	\$	T	F	Burger	>60	F
X_{10}	T	T	T	T	Full	\$\$\$	F	T	Italian	10-30	F
X_{11}	F	F	F	F	None	\$	F	F	Thai	0-10	F
X_{12}	T	T	T	T	Full	\$	F	F	Burger	30-60	T

What are the ACID transaction properties, and what can a database do to ensure each one?

Distinguish between the following:

- 1 a serial schedule
- 2 a serializable schedule
- 3 a conflict-serializable schedule

Is the following schedule conflict-serializable?

If not, what can the database transaction manager do to make the schedule conflict-serializable?

T1	r(x)	w(y)	
T2		r(y)	r(x)
T3		w(x)	r(x)